



Experiment 1: Quantization-Sparsity Pareto Exploration on IMG_DS


(a) Accuracy vs LUT




(b) Accuracy vs Warm E2E Latency



(c) Latency vs Sparsity



(d) Resource Utilization



All Ours configs <27 ms E2E → conservative >30× speedup vs ViT4Mal 0.81 s. Resource source: vivado_implementation. Accuracy on 3,444 test images, PYNQ-Z2.

30× (27 ms)